




Genetic algorithm

Python for AI

Optimizer: Gradient Descent

repeat until convergence {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$
 (for $j = 1$ and $j = 0$)
}



Optimization

- Derivative-based Optimization
 - Descent Methods
 - The Method of Steepest Descent
 - Classical Newton's Method
 - Step Size Determination
- Derivative-free Optimization
 - Genetic Algorithms
 - Simulated Annealing
 - Random Search
 - Downhill Simplex Search

Example: n -queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal
 - Note different search space... all states have N queens



- Is it a satisfaction problem or optimization?

8-queens problem


18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♛	13	16	13	16
♛	14	17	15	♛	14	16	16
17	♛	16	18	15	♛	15	♛
18	14	♛	15	15	14	♛	16
14	14	13	17	12	14	12	18

- Need heuristic function
 - Convert to an optimization problem
- h = number of ***pairs*** of queens attacking each other
- $h = 17$ for the above state

Search Space Recap


- State
 - All N queens on the board in some configuration
- Successor function
 - Move single queen to another square in same column.
- Example of a heuristic function $h(n)$:
 - the # of queens-pairs that are attacking each other
 - (we want to minimize this)

Hill climbing



Hill Climbing gets stuck in local maxima

Hill-climbing search: 8-queens problem



- Is this a solution?
- What is h ?
- Is any successor better?

Hill-climbing on 8-queens

- Randomly generated 8-queens starting states...
- 14% the time it solves the problem
- 86% of the time it get stuck at a local minimum
- However...
 - Takes only 4 steps on average when it succeeds
 - And 3 on average when it gets stuck
 - (for a state space with $8^8 \approx 17$ million states)


Hill-climbing with random restarts

- If at first you don't succeed, try, try again!
- Different variations
 - For each restart: run until termination vs. run for a fixed time
 - Run a fixed number of restarts or run indefinitely
- Analysis
 - Say each search has probability p of success
 - E.g., for 8-queens, $p = 0.14$ with no sideways moves
 - Expected number of restarts?

Use this algorithm!



Restarts	0	2	4	8	16	32	64
Success?	14%	36%	53%	74%	92%	99%	99.994%

- Expected number of steps taken?




Hey! Perhaps
sex can improve
search?

Sure! Check out
ye book.



Genetic algorithms


- Twist on Local Search: successor is generated by combining two parent states
- A state is represented as a string over a finite alphabet (e.g. binary)
 - 8-queens
 - State = position of 8 queens each in a column
- Start with k randomly generated states (**population**)
- Evaluation function (**fitness function**):
 - Higher values for better states.
 - Opposite to heuristic function, e.g., # non-attacking pairs in 8-queens
- Produce the next generation of states by “simulated evolution”
 - Random selection
 - Crossover
 - Random mutation



String representation
16257483

Can we evolve 8-queens through genetic algorithms?

Genetic algorithms



4 states for
8-queens
problem


2 pairs of 2 states
randomly selected based
on fitness. Random
crossover points selected

New states
after crossover

Random
mutation
applied

- Fitness function: number of non-attacking pairs of queens (min = 0, max = $8 \times 7/2 = 28$)
- $24/(24+23+20+11) = 31\%$
- $23/(24+23+20+11) = 29\%$ etc

Genetic algorithms



Has the effect of “jumping” to a completely different new part of the search space (quite non-local)

Example and code

- Download code in the classroom
- On class: follow a step by step tutorial

